

Fault tolerant TCP state

with Riak Core
Magnus and Luca

Use case

Problem: Some enterprise systems require keeping long-lived TCP connections

Solution: Separate TCP state from application and from the OS

- Enables application (or even OS!) restart without interrupting TCP connection

Problem: Storage of TCP state may be single point of failure

Solution: Replicate it - using Riak Core

Achievements

On TCP/IP handling:

- Design
 - One node per IP at a time
 - Minimal TCP segments handling
 - Received data synchronously passed to rest of application before storing updated state
- Sending raw IP packet

On replicated TCP state:

- Design

Apply a display filter ... <⌘/> Expression...

No.	Time	Source	Destination	Protocol	Dest port	Src pc
1	0.000000	127.0.0.1	127.0.0.1	IPv4		

- ▶ Frame 1: 40 bytes on wire (320 bits), 40 bytes captured (320 bits)
Raw packet data
- ▼ Internet Protocol Version 4, Src: localhost (127.0.0.1), Dst: localhost (127.0.0.1)
 - 0100 = Version: 4
 - 0101 = Header Length: 20 bytes (5)
 - ▶ Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
Total Length: 20
Identification: 0x0000 (0)
 - ▶ Flags: 0x00
Fragment offset: 0
Time to live: 64
Protocol: TCP (6)
 - ▼ Header checksum: 0x0000 [validation disabled]
[Good: False]
[Bad: False]
Source: localhost (127.0.0.1)
Destination: localhost (127.0.0.1)
[Source GeoIP: Unknown]
[Destination GeoIP: Unknown]

Software used

Riak Core - for replication of data

Tunctl - for sending and receiving raw IP packets (containing TCP segments)

Pkt - codec for IPv4 and TCP